

Data Science Engineer Resume

Boris Dev

May 22, 2025

Stack

- **AI:** Temporal, prompt engineering
- **Data Science:** Pandas, Sklearn, Numpy, Plotly, PySpark, DataBricks, Jupyter, SageMaker, PyTorch, GroundTruth, SQLAlchemy
- **Web BE:** FastAPI, Postgres, Mongo, Prometheus, Grafana, Splunk, ELK, OTEL
- **Cloud:** Azure, AWS
- **Pipelines:** Jenkins, Kafka

Education

PhD in Quantitative Human Geography, at SDSU and UCSB, 2015. Data science for location referenced social science problems. Dissertation: [New Metrics for Assessing Inequality using Geographic Data](#)

Experience

Sindri, 2025 - current, Consultant

Built Temporal AI Agent Evaluation framework

- AI system eval workflow: expectations yaml spec -> snapshot workflow -> evaluate and report
- LLM-as-judge: Prompt eval and fine-tuning workflow: faults spec -> prompt patch -> load -> batch find faults -> aggregate-score-enhance ** step 0: patch the AI system with a candidate prompt ** step 1: build a batch of input test examples synthetically ** step 2: write a “faults spec” markdown files explaining to the LLM what counts as a major and minor fault. ** step 3: Have the LLM list all minor and major faults for each example case ** step 4: Consolidate all the faults and have the LLM summarize them, score the prompt (0-5) given a guidance criteria, give a score rationale, and offer a short list of prompt changes to make a new candidate prompt, then repeat (iterate)
- Developers evaluate using a cli tool (python click library)
- SME-authored executable spec file, `sme-spec.yaml` [in progress]

Nobsmed, 2024 - current, Founder

Built LLM based website, <https://nobsmed.com/>, for users to search for biohacking experiences found in clinical studies and Reddit comments.

- ETL: Language AI parsing of 100,000 studies and 1,000,000 Reddit comments.
- Topic modeling: Extended Bertopic using LLM for naming and classification of outliers, [bertopic-easy](#).
- Customer discovery and user interviews.
- Experimented with fine-tuning opensource embeddings model and ML classification.
- Prompt engineering using InstructorAI.
- Embedding DBs for search (Weaviate, ElasticSearch).
- Worked closely with the CTO on requirements
- Built the story generation Flask API service

Smaller consulting gigs

- The Program Labs, 2025 - Building LLM AI automation tooling and a Google Ads Experimentation platform for a non-profit accelerator that helps startups in disadvantaged communities.
- EcoR1, 2025 - Built LLM AI scraper to extract earnings calls
- Wolf Games, 2023 - Built [new graph prompting for story generation](#)
- Intuitive Systems, 2023, NED (Name Entity Disambiguation) of vendor receipts fed into AMD's sales analytic pipeline. Used LangSmith for evaluation.

AI Engineer consultant at SimpleLegal, 2022-2023

SimpleLegal is a legal billing analytics company.

- Built the ground truth for a small language model (SLM) for an AI feature to flag invoice billing discrepancies
- Managed SMEs (subject matter experts) and 5 workers on the annotation pipeline of the training examples
- Built the inference server (SageMaker, PyTorch) and its product integration

Lead Analytic Endpoint Engineer at Sight Machine, 2018-2021

Sight Machine is a manufacturing analytics company.

- Built the backend engineering on biggest public facing analytic feature
- Coordinated QA process with sales and engineering
- Built company's first distributed tracing
- Built the containerized frontend engineering development environment

Lead Data Engineer at HiQ Labs, 2015-2018

HiQ Labs was a people analytics company.

- Built the scraping system and observability (Splunk)
- Led the migration from a monolith to a micro-service paradigm
- Migrated the data science team to DataBricks PySpark and microservices

Developer at Urban Mapping, 2011-2013

Urban Mapping provided geospatial analytics to Tableau.

- Formulated first map rendering and query performance metrics integrated into our CI/CD pipeline

Impactful projects

- Revived a stuck AI feature by shifting the team's focus from training data quantity to quality
- Reduced Tableau customer complaints by building a new observability system and CI/CD pre-commit metrics
- Reduced the data science team's firefighting by building a microservice architecture
- Built a gaming company's first murder mystery story generator by chaining to prompts to force consistency ([post](#)).

Papers and code

- LLM based taxonomy (topic modeling): [bertopic-easy](#).
- [Language AI Evaluation 101: Know your user](#)
- [Langchain PR: Causal Program-aided Language \(CPAL\)](#). See Harrison Chase's [Tweet](#).
- [Work papers](#)
- [Academic papers](#)

Non-tech fun points

- Climbed Cotopaxi (21,000 ft).
- Bodyboarded Mexpipe.
- Taught with students in Medellín, Columbia to make an open-source geo clustering library (ClusterPy)

- Taught kids snowboarding, as an instructor
- Managed service workers, as a restaurant assistant manager
- Counseled severely emotionally disturbed children