Boris Dev: Staff Data Scientist at Fastly

Below I address your job posting section, *"We'll be super impressed if you have experience in one or more of these".*

Domain expertise in our industry or related fields

- built a scraping system to get around LinkedIn's bot detection system, which might be related to Fastly tools to block scraping.

Experience deploying and maintaining production machine learning systems

- deployed AWS SageMaker inferences server based on several classifier models
- migrated from a monolith to a microservice ML pipeline.
- built an observability system on ML training example ingestion.

leading technical teams or mentoring junior data scientists

- trained data scientists in PySpark and DataBricks..
- trained data scientists to convert their monolith code into deployable microservices.
- led a team of 1 engineer, 4 annotators, 1 data scientist, and 2 subject matter experts

Building AI-powered products

- deployed an AI feature that had previously been stuck for a year before my arrival
- For Nobsmed.com I use an LLM to extract semantic meaning from studies and posts.

Experience with specialized data types

- formulated social science inequality metrics around geospatial data.

prompt engineering and RAG systems

- For a Wolf Games, a murder mystery gaming company, I built their first story generation service using a automated DAG prompting with generative AI.
- At Nobsmed.com, I use an LLM to extract semantic meaning from clinical studies and Reddit posts that are indexed into a RAG system.

Contributions

- Tweet by LangChain
- Langchain PR: Causal Program-aided Language (CPAL)
- bertopic-easy
- Language AI Evaluation 101: Know your user

AI services

- Azure's OpenAI services.