Boris Dev | boris.dev@gmail.com

Data leader background

- Led Language AI model development to GA release
- Led a data engineering switch from monolith to microservice paradigm
- Led building new observability systems
- Leader in the GA release of a large analytic product feature
- PhD research in creating new data analysis methods to track inequality
- Engineering tools: microservice backend stack, Splunk, Lightstep
- Data science tools: PySpark, Pandas, Numpy, BI (Splunk, Tableau)

Career goal

My goal is to serve a team of scientists that are discovering new analytic methods to help decision makers grow the business. My interests include the following.

- Reduce scientist toil by optimizing and automating data processes
- Document a ubiquitous language of customer meaning behind the data
- Get more feedback through observability, prototypes (demos), and papers
- Bring engineering best practices to scientists

Leadership style

1. Questioning - cognitive empathy to get context to execute the below points.
2. Visioning - win-win attitude that results in joyful effort over obfuscating
3. Coaching - top-down feedback, smaller chunks, metrics, short-cuts
4. Serving - bottom-up feedback, documentation discipline, blockers

Job experience

Technical AI Product Manager, consultant at SimpleLegal, 2022-2023

SimpleLegal software is a tool for big companies to manage bills from their vendor law firms. Before I arrived, my employer's new language AI feature for flagging suspicious invoice lines was stuck for one year because of bad performance, even after investing in more human annotation. My job was to increase AI performance. My actions resulted in the feature being launched and winning an AI Excellence Award.

Managed the Senior Data Scientist consultant

- The technical vision I set for the Data Scientist and I was to launch an MVP in six months by discovering "low-hanging fruit" that could scale back the machine learning effort. This resulted in scaling back our machine learning training from eleven to three redefined classifiers. The triage was to kill one classifier that was never needed, replace two with expert rules, and stop the labeling on five whose performance was sufficient.

- To increase clarity, I coached the team to understand the customer meaning behind the data and the AI flags.
- To increase clarity, I crunched the knowledge of the Data Scientist's metrics analysis into a working paper, which was continually shared at the weekly stand-ups.
- I wrote the data scientist a stellar performance review, highlighting his acumen at articulating complex concepts to executives.

Managed the engineering tasks

- I assigned Jira issues to engineers.
- 1:1 meetings with the VP of Engineering.
- After one engineer left for another company, I stepped in to refactor the inference server.

Managed the offshore annotators

- Included the offshore annotators and subject-matter expert into a continuous improvement collaboration on AI flag definition and annotation guidelines.
- Changed annotation process to small batches
- Started inter-annotator performance metrics and meetings.
- Labeling several hundred sentences, I built cognitive empathy to identify convoluted annotation guidelines as the root cause of our poor labeling quality.

I changed our direction from increasing data quantity to increasing quality.

- I built relationships with five experts in three sister companies across the conglomerate.
- I included the subject-matter expert into our team.
- I learned from the annotation expert that success depended on including the annotators as experts for feedback.
- I nudged the data scientists to articulate the biggest issue as *"garbage in, garbage out"*.
- I connected with a Founder (VP of Product) at a sister company who gave me the pre-processing noise filter idea.
- I asked the President to convene a Amazon 6-pager meeting to share the new knowledge.

Lead Analytic Endpoint Engineer at Sight Machine, 2018-2021

The Sight Machine product is a performance dashboard tool for process engineers at factories.

- got the company's biggest public-facing feature, Recipes, across the GA line
- worked across silos (Customer Success, DevOps, Data Engineers)

- wrote high-level design papers to help the Product and Customer Success teams get relevant customer feedback on product technical design decisions.
- started a new checklist process to lock down successful sales demos.
- nudging the engineering delivery team to be closer to a product team, the Director of Engineering and I started a new process where the engineers fleshed out their own Jira issues. The new technical design autonomy reduced engineering toil and increased creativity.
- built the first distributed tracing (Lightstep), which made debugging system downtime easier.
- proposed and standardized frontend development environment, which made on-boarding easier.
- coached junior engineers.

Lead Data Engineer at HiQ Labs, 2015-2018

The HiQ product predicted the likelihood that an employee would quit.

- A puzzle for the CTO and I was to figure out how to get around LinkedIn's bot detection in order to scrape millions of HTML public profiles, the raw data for our prediction pipeline.
- I led a Junior DevOps engineer to build a Splunk observability system that was used to track performance and experiments on different spider configurations.
- I trained a Junior Data Engineer to maintain the scraping system.
- I led the move from a monolith to a microservice paradigm.
- I led a Data Engineer to guide the Data Scientists into a new microservice development paradigm.
- I explained to management the technical and human vision behind the microservice architecture.
- I shadowed the Data Scientists to see what I could automate for them.
- I migrated the Data Scientists from Mongo queries to PySpark.

Start-up partner and full-stack developer at Map Decisions, 2014

Map Decisions was a mobile app to automate street sign inspection for city Public Works departments.

Start-up partner and analytics developer at Geoscore, 2014

Geoscore (repo) ranks and visualizes 100,000+ neighborhoods.

Developer at Urban Mapping, 2011-2013

Urban Mapping provided a geodata REST API to Tableau and other customers.

- I built the observability system (Splunk).
- I created a new performance metric for map tile rendering.
- I built the CI pipeline (Jenkins).

- I built a Tableau BI Dashboard

Misc leadership roles

- coached students in Medellín, Columbia to make ClusterPy
- coached graduate students and taught spatial statistics, as a PhD
- Kids snowboard instructor at Vail Resorts, CO.
- Restaurant Assistant Manager at Vail Resorts, CO.
- Lead analyst at California Business Magazine. CA. Collected company data for newspaper rankings.
- Counselor for emotionally disturbed children. Seneca Institute. CA.

Open source code and writing

- A play Ethereum MEV bot
- A git bare approach to version control your dot files
- *Why did your language AI feature fail?*
- Work papers
- PhD thesis. *Assessing Inequality using Geographic Income Distributions*
- *Spatial Econometrics* entry. Encyclopedia of Human Geography. 2009.
- *Interactive spatio-temporal modelling of health systems*
- *σ-convergence in the presence of spatial effects*
- *Integrating Econometric and Input-Output Models in a Multiregional Context*