Boris Dev | boris.dev@gmail.com

Interests

- Process optimization to scale humans
- Documenting a ubiquitous language to scale humans
- Building observability systems and prototypes for quick feedback

Tools

- **Backend engineering:** Linux, Python, Postgres, Flask, Django
- **DevOps:** Docker, Jenkins, Helm Charts, AWS Sagemaker and Lambda
- **Observability:** Splunk, Lightstep for distributed tracing
- **Data science:** SQLAlchemy, PySpark, Pandas, Numpy, PyTorch, Tableau
- **Teamwork:** Inviting feedback with demos and papers (examples).

Job experience

Data Scientist / Technical AI Product Manager, consultant at SimpleLegal, 2022-2023

SimpleLegal software is a tool for legal departments of big companies to manage bills from their vendor law firms. Their problem was that their new language feature for flagging suspicious invoice lines was stuck for one year because of bad performance, even after investing in more human annotation. My job was to increase AI performance. The actions I took resulted in the company launching their first AI feature, with positive feedback from customers and the sales team, as well as an AI Excellence Award.

Managed the Senior Data Scientist consultant

- The technical vision I set for the Data Scientist and I was to continually search for "low-hanging fruit" to scale back machine learning to launch an MVP. This resulted in scaling back our machine learning training from eleven to three classifiers. The triage was to kill one classifier that was never needed, replace two with expert rules, and stop the labeling on five whose performance was sufficient.
- To increase clarity, I crunched the knowledge of the Data Scientist's metrics analysis into a working paper, which was continually shared at the weekly stand-ups.
- Periodically, I asked the Data Scientist to take the lead in explaining complex concepts to the rest of the team, including executives. I wrote him a stellar performance review, which highlighted his capacity to articulate complex concepts to executives.

Managed the inference server engineering and QA tasks

- I assigned Jira issues to engineers.
- 1:1 meetings with the VP of Engineering.

- After one engineer left for another company, I stepped in to refactor the inference server.
- I organized the subject-matter expert embarrassment reviews.

Managed the offshore annotators

- I overhauled the relationship with our team of offshore annotators from a hierarchical one to a democratic one.
- To understand the perspective of the offshore annotators, I labeled several hundred sentences. The result was I identified convoluted annotation guidelines as the root cause of our poor labeling quality.
- I asked the manager of the annotators to let his workers interact directly with me and the subject-mater through Slack to give us challenging feedback on our annotation guidelines (the shared language document).
- I developed and shared new performance metrics with the team over periodic Zoom meetings.

Built relationships across the conglomerate

- I built relationships with experts spread across three sister companies: a principle Data Scientist, two subject-matter experts (one of which was a VP of Product/Founder), a DevOps engineer, an MLOps engineer, a machine learning expert, and a human annotation expert.
- From questioning the subject-matter expert, I figured out we were missing pre-processing noise filters.
- The annotation expert helped me recognize that we needed a democratic relationship with the annotators to ensure quality.
- After gathering the evidence, the Data Scientist, MLOps engineer, and I reached a consensus: the biggest potential factor to bad performance was *"garbage in, garbage out"*. I stopped the human annotation process. I asked the company's President to convene a meeting with our expanded team in the style of the Amazon 6-pager. The attendees spent the first fifteen minutes reading my high-level design document to ensure we were on the same page with complex issues before the discussion started.

Backend developer at Sight Machine, 2018-2021

The Sight Machine product is a dashboard tool for process engineers at manufacturing plants to keep track of performance.

I built relationships across silos (Customer Success, DevOps, Data Engineers, and Product Engineers).

- To bring the company's biggest public-facing feature, Recipes, from its embryonic start as a spreadsheet to general release, I shared high-level design papers and gave non-technical explanations to the Product and Customer Success teams. This allowed them to get accurate customer feedback on product technical design decisions.

- I started a new process to ensure toil free and quality sales demos. Engineers must be included in the calendar invites. Twenty-four hours before the scheduled demo, the environment's compute workers are scaled up, the best frontend and backend branches deployed, and lastly, the product manager starts smoke and sanity testing.

I led the following initiatives to increase the engineering team's productivity.

- The Director of Engineering and I started a new process where the engineers wrote their own Jira issues, instead of being delegated Jira issues by the Product team. The new technical design autonomy allowed us to meet requirements with cheaper technical implementations.
- To debug faster, I instrumented and built distributed tracing (Lightstep) and managed the DevOps work.
- To simplify on-boarding, I containerized the frontend development environment.
- I coached junior engineers.

Data and product engineer tech lead at HiQ Labs, 2015-2018

Our product provided predictions to our customers on whether their employees were about to quit.

Tech lead for scraping

- The puzzle for the CTO and I was to figure out how to get around LinkedIn's bot detection in order to scrape millions of HTML public profiles, the raw data for our prediction pipeline.
- I led a Junior DevOps engineer to build a Splunk observability system that was used to track performance and experiments on different spider configurations.
- I trained a Junior Data Engineer to maintain the scraping system.

Tech lead to move the company from a monolith to a microservice architecture

- I coached a Data Engineer to guide the Data Scientists into a new microservice development pattern.
- I explained to management the technical and human vision behind the microservice architecture.

Supported Data Scientists to reduce their toil

- I shadowed the Data Scientists to see what I could automate for them.
- I migrated the Data Scientists from Mongo queries to PySpark.
- To scale, I trained a Junior Data Engineer to train the Data Scientists on the new development environment that I had set up.

Start-up partner and developer at Map Decisions, 2014

We provided a mobile app to automate street sign inspection for the Public Works departments of cities. I built the app.

Developer at Urban Mapping, 2011-2013

We provided a location query and map tiling service to Tableau. I built the company's first observability system and QA system.

Misc jobs

- As a PhD researcher, I coached master's degree program students.
- As a PhD researcher, I was a teaching assistant in a spatial statistics class.
- Kids snowboard instructor at Vail Resorts, CO.
- Assistant Manager at Gundy's Grill restaurant at Vail Resorts, CO.
- Researcher on building a prototype of the Wall Street Journal's "Main Street" section. California Business Magazine. CA.
- Led analysts collecting regional company ranking for large newspapers. California Business Magazine. CA.
- Counselor for emotionally disturbed children. Seneca Institute. CA.

Writing

- Why did your language AI feature fail?
- Work papers
- PhD thesis. Assessing Inequality using Geographic Income Distributions
- Spatial Econometrics entry. Encyclopedia of Human Geography. 2009.
- Interactive spatio-temporal modelling of health systems
- σ-convergence in the presence of spatial effects
- Integrating Econometric and Input-Output Models in a Multiregional Context

Open source code

- A play Ethereum MEV bot
- A git bare approach to version control your dot files
- Geoscore website and repo to rank and visualize 100,000+ neighborhoods.
- ClusterPy. As co-founder, I coached students in Medellín, Columbia in geographic data software.