

Boris Dev

Staff Software Engineer, AI Platform

San Francisco • boris.dev@gmail.com • [github](#) • [linkedin](#)

AI Product • AI Evals • Analytics • Data Engineering • Geospatial (PhD Quantitative Geography)

Experience

NoBSmed, 2024 - Present, Founder

Building a causal-evidence diff engine that compares a patient's current health approach against clinical-study evidence, surfacing omitted upstream drivers, downstream effects, contested evidence, and questions to investigate with a doctor or AI.

- Built the causal-graph diff engine as a deterministic DAG dataflow with LLM calls inside individual nodes: Patient Context → Evidence Retrieval → Structured Findings → Evidence Graph → Diff/Evaluation → Report.
- Built hybrid retrieval (RAG) over Azure AI Search and a Neo4j graph database with a SNOMED ontology layer.
- Designed provenance-linked, ontology-grounded, typed representations of patient context and clinical-study findings — comparable, queryable, and usable for causal reasoning and downstream AI evaluation.
- Built incremental evidence-graph construction with reusable parsed findings cached in Databricks Delta.
- Ran Reddit GTM experiments identifying demand for evidence-based audits of medication and supplement stacks, particularly among biohacking communities.

Sindri, Oct 2025 - Feb 2026, Consultant

Sindri is an early-stage startup applying AI to document management for large energy-industry construction projects.

Built the team's first AI evaluation framework, replacing engineer-driven manual QC/QA with automated checks and unblocking high-stakes customer demos by lifting AI email quality.

- Designed an SME-authored YAML expectations DSL (pre-run scenarios + post-run predicates) so domain experts — not just engineers — could specify what “correct” looks like for a Temporal workflow run
- Built a Temporal-aware test harness that snapshots post-run database side effects and activity outputs, then evaluates each expectation — became the team's foundational CI/CD for iterating on Temporal modules
- Built an LLM-as-judge pipeline for the AI task of writing customer emails that explain how to fix their supplier non-conformance cases — scoring candidate prompts against synthetic test batches and emitting a structured fault taxonomy (top faults, rationale, proposed prompt edits) to drive iteration

Smaller consulting gigs

- EcoR1, 2025 - LLM extraction of earning-call calendar events
- Intuitive Systems, 2023 - LLM extraction of AMD products from vendor receipts; LangSmith for evaluation

AI Engineer consultant at Wolf Games, 2023-2024

Wolf Games is a murder-mystery gaming company piloted by the producers of Law & Order.

- Fixed story generation to be consistent by building a DAG-based story-composition engine that dynamically chained LLM prompts to maintain narrative coherence across overlapping multi-step workflows, ensuring consistency in plot and in character MMOs (Means, Motive, Opportunity). [Read Google AI showcase here](#)

AI Engineer consultant at SimpleLegal, 2022-2023

SimpleLegal is a legal billing analytics company.

- Identified a poorly specified rubric as the root cause of low model quality on a stuck feature
- Designed a collaborative process for paralegals and lawyers to debate edge cases, build consensus, and elicit the nuanced expertise needed to refactor the rubric
- Built a quality-control annotation pipeline around the new rubric → massive increase in training-example quality and the launch of the previously stuck feature
- Deployed a PyTorch Small Language Model on SageMaker and the ML client into the Flask product app

Lead Analytic Endpoint Engineer at Sight Machine, 2018-2021

Sight Machine is a manufacturing analytics company.

- Built the backend engineering on the biggest public-facing analytic feature
- Implemented a pre-demo protocol between product and engineering → less panic before each sales demo
- Coordinated QA process with sales and engineering → better prioritization/triage
- Built the company's first distributed tracing → simpler firefighting for mid-level developers
- Containerized the frontend build → standardized the team's setup & scaled testing to cloud

Lead Data Engineer at HiQ Labs, 2015-2018

HiQ Labs was a people analytics company.

- Taught data scientists how to refactor their pipeline code into microservices
- Refactored the scraping system → established pipeline reliability
- Refactored the data pipeline from a data-science monolith to a microservice paradigm → established release reliability
- Migrated the data science team from Mongo to PySpark/Databricks → increased productivity on new-product R&D

Developer at Urban Mapping, 2011-2013

Urban Mapping provided geospatial analytics to Tableau.

- Built the first observability for the geospatial API (delivered to Tableau) → surfaced system-performance metrics to prioritize coding issues
- Built the first performance-regression gate → reduced failed releases/customer complaints

Small startup attempts

- **MapDecision** — geospatial map client for Public Works departments
- **GeoScores** — geospatial analysis tool for neighborhood quality

Recent open-source projects

- **tau-discernment** — grades agent *discernment* in balancing competing goals (*in progress*).
- **nobsmed-healthbench-audit** — flagged 29 decision-changing errors in OpenAI's HealthBench.
- **healthbench-ebm-verified** — failure-pattern analysis of GPT-5.2 & Claude Opus 4.8 on medical AI questions.

Stack

Core: Python · Pydantic (heavy) · FastAPI · Django / DRF · PydanticAI (light) · LangGraph (light) · Neo4j · Postgres · MongoDB · BM25 / Azure AI Search · Databricks / Delta / PySpark · Temporal (early) · AWS · Azure · OpenTelemetry · TypeScript (light) · Angular (older)

Education

PhD in Quantitative Human Geography at SDSU and UCSB, 2015. Data science for location-referenced social-science problems. Dissertation: [New Metrics for Assessing Inequality using Geographic Data](#)

Papers & code	Non-tech fun
LLM-based taxonomy (topic modeling): bertopic-easy	Climbed Cotopaxi (21,000 ft)
Language AI Evaluation 101: Know your user	Bodyboarded Mexpipe
Langchain PR: Causal Program-aided Language (CPAL) — see Harrison Chase's tweet	Taught coding with students in Medellín, Colombia
ClusterPy — open-source geo clustering library	Taught kids snowboarding as an instructor
Work papers	
Academic papers	Counseled severely emotionally disturbed children